

ARTIFICIAL INTELLIGENCE

NATURAL DISASTER INTENSITY ANALYSIS AND CLASSIFICATION

ASSIGNMENT 4

READ DATASET AND DO PREPROCESSING

Input

```
df = pd.read_csv(r'spam.csv',encoding='latin-1')
df.head()
```

Output

	v1	v2	unnamed:2	unnamed:3	unnamed:4
0	ham	Go until jurong point, crazy.. Available only ...	NaN	NaN	NaN
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3	ham	U dun say so early hor... U c already then say...	NaN	NaN	NaN
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN	NaN	NaN

Input

```
df.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed:4'],axis=1,inplace=True)
df.info
```

Output

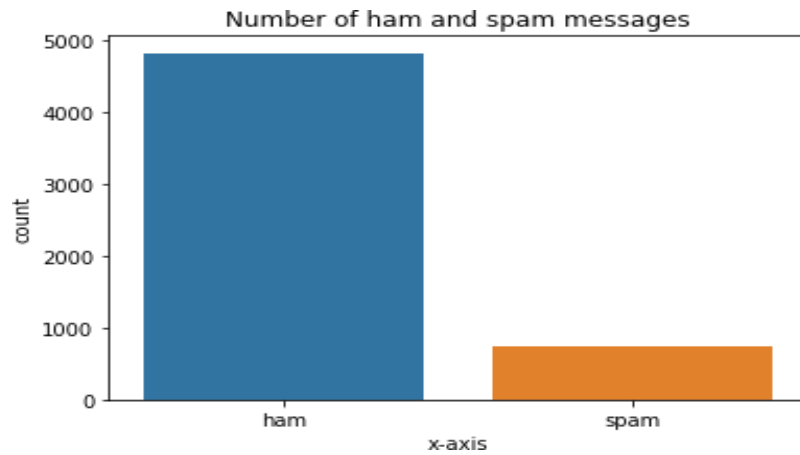
```
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0    v1      5572 non-null   object
 1    v2      5572 non-null   object
dtypes: object(2)
memory usage: 87.2+ KB
```

Input

```
sns.countplot(df.v1)
plt.xlabel('x-axis')
plt.title('Number of ham and spam messages')
```

Output

```
Text(0.5, 1.0, 'Number of ham and spam messages')
```



CREATE INPUT VECTORS AND PROCESS LABELS

```
X = df.v2
```

```
Y = df.v1
```

```
le = LabelEncoder()  
Y = le.fit_transform(Y)
```

```
Y = Y.reshape(-1,1)
```

SPLIT THE TRAINING AND TESTING DATA

```
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.20)
```

PROCESS THE DATA

```
max_words = 1000  
max_len = 150
```

```
tok = Tokenizer(num_words=max_words)  
tok.fit_on_texts(X_train)
```

```
sequences = tok.texts_to_sequences(X_train)  
sequences_matrix = sequence.pad_sequences(sequences,maxlen=max_len)
```

CREATE MODELS AND ADD LAYERS

Input

def RNN():

```
inputs = Input(name='inputs', shape=[max_len])
layer = Embedding(max_words, 50, input_length=max_len)(inputs)
layer = LSTM(128)(layer)
layer = Dense(256, name='FC1')(layer)
layer = Activation('relu')(layer)
layer = Dropout(0.5)(layer)
layer = Dense(1, name='out_layer')(layer)
layer = Activation('tanh')(layer)
model = Model(inputs=inputs, outputs=layer)
return model
```

```
model = RNN()
```

```
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
inputs (InputLayer)	[(None, 150)]	0
embedding (Embedding)	(None, 150, 50)	50000
lstm (LSTM)	(None, 128)	91648
FC1 (Dense)	(None, 256)	33024
activation (Activation)	(None, 256)	0
dropout (Dropout)	(None, 256)	0
out_layer (Dense)	(None, 1)	257
activation_1 (Activation)	(None, 1)	0
=====		
Total params: 174,929		
Trainable params: 174,929		
Non-trainable params: 0		

```
model.compile(loss='binary_crossentropy',optimizer=RMSprop(),metrics=['accuracy', 'mse', 'mae'])
```

FIT THE MODEL

Input

```
model.fit(sequences_matrix,Y_train,batch_size=128,epochs=100,
```

Output

```
validation_split=0.2,callbacks=[EarlyStopping(monitor='val_loss',min_delta=0.0001)])
```

```
Epoch 1/100
```

```
28/28 [=====] - 13s 460ms/step - loss: 0.0961 - accuracy: 0.9778 - mse: 0.0358 - mae: 0.1438 - val_loss: 0.1271 - val_accuracy: 0.9832 - val_mse: 0.0568 - val_mae: 0.2060
```

```
Epoch 2/100
```

```
28/28 [=====] - 14s 507ms/step - loss: 0.0728 - accuracy: 0.9885 - mse: 0.0607 - mae: 0.2129 - val_loss: 0.1175 - val_accuracy: 0.9821 - val_mse: 0.0766 - val_mae: 0.2416
```

Input

```
test_sequences = tok.texts_to_sequences(X_test)
test_sequences_matrix = sequence.pad_sequences(test_sequences,maxlen=max_len)
```

```
accr = model.evaluate(test_sequences_matrix,Y_test)
```

Output

```
35/35 [=====] - 3s 92ms/step - loss: 0.1390 - accuracy: 0.9821 - mse: 0.0779 - mae: 0.2393
```

Input

```
print('Test set\n Loss: {:.3f}\n Accuracy: {:.3f}'.format(accr[0],accr[1]))
```

Output

```
Test set
  Loss: 0.139
Accuracy: 0.982
```

SAVE THE MODEL

```
model.save(r"C:\Users\aruna\OneDrive\Desktop\model_lstm.h5")
```

TEST THE MODEL

Input

```
from tensorflow.keras.models import load_model
m2 = load_model(r"C:\Users\aruna\OneDrive\Desktop\model_lstm.h5")

m2.evaluate(test_sequences_matrix,Y_test)
```

Output

```
35/35 [=====] - 4s 68ms/step - loss: 0.1390 -
accuracy: 0.9821 - mse: 0.0779 - mae: 0.2393
```

```
[0.13899557292461395,
 0.9820627570152283,
 0.07788368314504623,
 0.23931345343589783]
```

